How to Use Qyro

> *Qyro generates unique, high-quality Q&A sets on any CS topic using AI. This guide walks you through everything from first visit to downloading your questions.*

Table of Contents

1. First Visit — Getting Started
2. Step 1 — Add Your API Keys
3. Step 2 — Generate Questions
4. Step 3 — Read Your Results
5. Chat History
6. Tips & Best Practices
7. FAQ

---

1. First Visit — Getting Started

Open the site at **https://qyro-ten.vercel.app**.

- On first load, Qyro automatically creates an **anonymous account** for you using a UUID stored in your browser. You don't need to sign up or log in.
- Your session and history are tied to this browser. **Don't clear your browser storage** or you'll lose your history.

> *[!IMPORTANT] The site requires **your own API keys** to generate questions (see Step 1). Nothing works without them.*

---

2. Step 1 — Add Your API Keys

Before generating anything, you must add your API keys. Qyro uses a **"bring your own keys"** model — your keys are stored only in the app's database under your anonymous ID, never shared.

How to open Settings

Click the ⚙ **gear icon** in the top-right corner of the interface.

Keys you need

Key	Where to get it	Required?
Gemini API Key	aistudio.google.com → Get API key	✅ Yes
Gemini API Key 2	Same as above — a second key to split load	Optional
Groq API Key	console.groq.com → API Keys	✅ Yes
Groq API Key 2	Same, a second Groq key	Optional
HuggingFace Token	huggingface.co/settings/tokens → New token (read)	✅ Yes
Custom Database URI	A hosted PostgreSQL string (e.g. from Neon)	Optional

Steps

1. Paste each key into its field.
2. Click **Save**.
3. Close the settings panel.

> *[!TIP] **Why two keys for Gemini/Groq?** The agent runs multiple parallel pipelines simultaneously. Having two keys prevents hitting rate limits during big generation runs.*

> *[!NOTE] **Custom Database URI** — If you provide your own hosted PostgreSQL URI (like from [Neon](#) or Supabase), Qyro saves generated Q&As directly to your database. `localhost` **strings will NOT work** because the cloud app cannot reach your personal computer. If left blank, questions are saved to Qyro's shared database.*

---

3. Step 2 — Generate Questions

Once your keys are saved, use the form in the main chat area to configure your generation run.

Configuration fields

Subject

Pick from a categorized list of CS topics:

Category	Example Subjects
Core CS	DSA, DBMS, Operating Systems, Computer Networks
Software Engineering	OOP, System Design, Software Engineering
Web & Mobile	Web Development, Mobile App Development
AI & Data	Machine Learning, Data Science, AI
Infrastructure	DevOps, Cloud Computing
Security & Web3	Cybersecurity, Blockchain

Difficulty

Choose how hard the questions should be:

- **Easy** — Recall and basic concept questions
- **Medium** — Requires understanding of internals and trade-offs
- **Hard** — Deep analysis, edge cases, and production-level reasoning

Bloom's Taxonomy Level

Controls the *cognitive depth* of questions. Leave blank for a mix.

Level	What it tests
L1 — Recall	Can you remember the fact?
L2 — Understand	Can you explain it?
L3 — Apply	Can you use it in a scenario?

L4 — Analyze	Can you break it down and compare?
L5 — Evaluate	Can you judge and justify?
L6 — Create	Can you design something using it?
L7 — Innovate	Can you propose something new?
Mixed Bloom	Automatically distributes across levels

Number of Questions (N)

How many questions to generate **per parallel run**. Start with `5` if you're testing.

Parallel Workflows

How many independent generation pipelines to run simultaneously. `2` parallel workflows with `N=5` will attempt to produce up to 10 questions (duplicates are automatically removed).

> *[!TIP] **Recommendation:** Start with `N=5`, `Parallel Workflows=2`. Increase once you're familiar with generation time.*

Start generation

Click **Generate** (or press Enter). The button will be disabled while generation is in progress.

---

4. Step 3 — Read Your Results

Generation happens in real time — you'll see status messages appear in the chat as the agent works:

```
Starting generation...
Generating questions...
Checking for duplicates...
Generating answers...
Done! X questions generated.
```

Question cards

Each generated question appears as a card with:

Field	Description
Question	The full question text
Options (A/B/C/D)	Multiple choice options
Correct Answer	The right option, highlighted
Answer Explanation	A detailed, dense explanation
Bloom Level	The cognitive level this question targets
Difficulty	Easy / Medium / Hard
Topic Tags	Specific sub-topics this question covers

Click on a card to expand the full answer and explanation.

---

5. Chat History

Every generation run is saved as a **chat session** in the left sidebar, labeled with the subject you chose.

Navigating history

- Click any session in the sidebar to reload that conversation.
- Sessions are ordered **newest first**.
- Each session shows the subject name and how many questions were generated (`Q: 5`).

> 💡 *History is tied to your anonymous browser session. It persists across page refreshes as long as you don't clear your browser's local storage.*

---

6. Tips & Best Practices

Getting the best questions

- **Be specific with subject** — "Data Structures & Algorithms" will generate more focused questions than a broad topic.
- **Use a subject description** — If the form has a description field, add context like *"Focus on graph algorithms and dynamic programming"* to guide the LLM.
- **Match difficulty to your goal** — Use `easy` to review fundamentals, `hard` to prepare for technical interviews.
- **Use Mixed Bloom for comprehensive sets** — A Mixed Bloom run produces variety from recall (L1) up to creative (L6/L7).

Rate limiting

- You can only start a new generation **once every 2 minutes** (per IP). This is intentional to prevent API key exhaustion.
- If you see a *"Please wait X seconds"* message, just wait before trying again.

Using two API keys

- If generation is slow or fails with rate limit errors, go to ⚙ Settings and add a **second Gemini key** and a **second Groq key**.
- The agent automatically alternates between key 1 and key 2 across parallel runs.

---

7. FAQ

Q: Do I need to create an account? No. Qyro uses anonymous UUID-based sessions. Your identity is stored in your browser.

Q: Where are my API keys stored? In Qyro's PostgreSQL database, linked to your anonymous UUID. They are never logged, printed, or shared.

Q: Why does generation take 30–90 seconds? The agent runs multiple LLM calls in parallel (question generation, duplicate checking, answer generation) and waits for all of them to complete before returning results. This is normal.

Q: I got fewer questions than I asked for — why? The deduplication pipeline removes questions that are too similar to ones already in the database OR similar to each other within the same run. This is by design — quality over quantity.

Q: Can I use `localhost` as my database? Only if you are running the **Agent (Backend) locally** on your computer. If your Agent is hosted on **Railway**, it cannot see your `localhost` (it will try to connect to itself on Railway's server). For the Cloud Agent, you must use a hosted database like **Neon**, **Supabase**, or **Railway's PostgreSQL**.

Q: Can I use Qyro for subjects other than CS? The subject list is currently fixed to CS topics. Custom subject input may be added in future versions.